

Group- and Time-Split Coverage in Tabular Models Fails to Certify, but Uncertainty Abstention Robustly Repairs Selective Risk: A Preregistered Audit with Held-Out Confirmatory Replication

Zeyu Fu*

July 17, 2026

Abstract

Uncertainty-ranked abstention beats random deferral on 13 of 14 public OpenML datasets with a defensible entity/time key, across four open, token-free configurations in three architecture families—XGBoost, LightGBM, and tabular foundation models TabICLv2 and TabDPT—surviving Holm–Bonferroni for the two gradient-boosted models (with and without the group key) and bootstrap CIs excluding zero for the two foundation models; a preregistered, held-out confirmatory replication on TableShift (Gardner et al., 2023) (9 credential-free tasks, 6 evaluable; frozen endpoint rule, $K=20$ split-repeats) confirms without exception that this free uncertainty score ties-or-beats a 12-trial-tuned learned selector (H2 CONFIRM). Beating random deferral is a low bar that any score with informative confidence clears; the substantive claims are that the advantage holds under grouped and temporal shift, that no tabular foundation model auto-repairs it, and that the free score matches a tuned one at no extra cost—though both sit far below an oracle ceiling (median repair ratio 0.339 free vs. 0.338 tuned vs. 0.745 oracle), leaving substantial unclaimed headroom.

This repair does not restore coverage. Tabular models are routinely validated with random splits, but deployment shifts are usually *grouped*—a new entity, region, or period—and moving to a leave-group-out / rolling-origin split leaves split-conformal coverage below its 0.90 nominal target: median worst-group coverage ≈ 0.81 on the exploratory 14-dataset audit and 0.797 on the confirmatory TableShift roster (6 evaluable). No tested distribution-free remedy (weighted or inferred-group conformal) restores it at the majority level, on either the exploratory audit or the preregistered held-out replication (H1 REJECT), so the negative-coverage/positive-selection prescription is not an artifact of the 14-dataset convenience suite. We recommend plain split-conformal plus uncertainty abstention—no foundation model, token, or Mondrian add-on—for grouped tabular shift, scoped strictly to *selective* reliability rather than a coverage guarantee.

Our own preregistered bars were majority counts: the repair count passed robustly, but the grouped-gap count that would certify a directional *collapse* is underpowered and knife-edge (its verdict flips on one split re-draw), so our conjunctive Stage-2 gate evaluated to KILL and we read the degradation as a direction rather than a robust, preregistered count.

1 Introduction

Tabular models are almost universally evaluated with random hold-out splits, yet the shifts that matter in deployment are almost never random: the fraud-detection model trained on known ac-

*ORCID: 0009-0001-8329-0108. Correspondence: fuzeyu09@gmail.com. State Key Laboratory of Trauma and Chemical Poisoning, Institute of Combined Injury, Army Medical University, Chongqing, China.

counts is applied to a new merchant; the demand model trained on historical months must forecast an upcoming quarter; the healthcare risk model trained at established sites is deployed at a new facility. The split that governs reliability at inference is a grouped or temporal split, not a random one. If random-split validation systematically overstates reliability under such shifts, then published performance figures are misleading and practitioners are blind to actual deployment risk—yet this gap is rarely closed with controlled experiments on standard benchmarks.

We ask three empirical questions about the primary phenomenon, plus one negative-control question about a candidate remedy. (i) Do grouped/temporal splits cause a measurable reliability collapse—degraded risk-coverage and conformal under-coverage? (ii) Does uncertainty-ranked abstention improve selective reliability over random deferral? (iii) Is the effect model-agnostic, reproducing across independent model families? As a negative control on a natural remedy candidate, we additionally ask (iv) whether an *uncertainty-taxonomy* (confidence-conditional) Mondrian conformal predictor helps; it does not (section 3), and we report that compact negative result and its own key-dependent explanation as a control rather than a co-equal finding. (We use the uncertainty-decile Mondrian because a group-keyed taxonomy provably degenerates to split conformal under leave-group-out, where the grouping key is absent from calibration.) Uncertainty-ranked abstention is an instance of selective classification (Geifman and El-Yaniv, 2017)—the learning-theoretic descendant of the classical reject option, from Chow (1970)’s error-reject tradeoff through statistically grounded classification with a reject option (Herbei and Wegkamp, 2006) to margin-based learning with rejection (Cortes et al., 2016)—and selective mechanisms are known to interact non-trivially with subgroup structure (Jones et al., 2021), with reject options themselves able to shape subgroup disparities (Kamiran et al., 2018)—a point our Folktables/ACS generalization arm returns to; the conformal machinery we audit is standard split (inductive) conformal (Papadopoulos et al., 2002; Vovk et al., 2005; Angelopoulos and Bates, 2023). We answer these questions on 14 public, keyed OpenML datasets with defensible entity- or time-split keys (Vanschoren et al., 2014), using two open gradient-boosted models (XGBoost (Chen and Guestrin, 2016), LightGBM (Ke et al., 2017)) and two open-weights tabular foundation models (TabICLv2 (Qu et al., 2026), TabDPT (Ma et al., 2024)), all token-free.

Two results are the ones we hold ourselves to a preregistered, held-out standard for, and both hold. First, uncertainty-ranked abstention beats random deferral on 13 of 14 datasets for each of four open, token-free configurations spanning three architecture families—two gradient-boosted models (XGBoost, LightGBM), which behave near-identically, plus two distinct tabular-FM architectures (TabICLv2, TabDPT)—with and without the group key; in the GBM arms it further holds across default and tuned hyperparameters (12–13/14) and in every one of $K=20$ –30 repeated split draws. Beating random deferral is, on its own, a low bar: any score whose confidence is even weakly informative about error will clear it, so the count by itself is not the substantive claim. What is substantive, and what a preregistered, held-out confirmatory replication on TableShift now confirms without exception at every split-repeat, is that the free uncertainty score ties-or-beats a budget-matched tuned selector (H2 CONFIRM; section 3) on data that could not have generated the hypothesis—so the advantage holds under grouped and temporal shift, is model-agnostic across three architecture families including two foundation models that do not auto-repair the underlying gap, and costs nothing over the model’s own confidence—though both the free and tuned scores remain well below an oracle ceiling (section 3: median repair ratio 0.339 free vs. 0.745 oracle), leaving room the field has not yet claimed. Second, the same held-out replication REJECTs coverage restoration (H1): no tested distribution-free remedy restores group-conditional coverage on a majority of evaluable tasks, corroborating rather than contradicting the diagnosis, since we never recommend those remedies as a coverage fix. What abstention repairs is the retained-set selective risk—a measured median $\sim 33\%$ reduction relative to random deferral (section 3)—not the cov-

erage gap itself: we do not claim to close a measured fraction of the random→grouped coverage gap (which we cannot, lacking persisted matched-coverage risk–coverage curves), to beat any baseline stronger than random deferral, or that abstention restores the coverage that degrades under grouped/temporal shift; these are different quantities, kept distinct throughout.

The degradation these two confirmed results are read against is real in direction—grouped splits degrade risk–coverage across all four configurations and push split-conformal coverage below the 0.90 nominal target on the temporal splits, including two tabular FMs that do not auto-correct the gap—but its dataset-count is not itself a robust, preregistered claim, and we say so plainly rather than let the positive results above imply otherwise: our conjunctive Stage-2 gate on that count evaluated to KILL, the confirmatory/exploratory partition we draw throughout is accordingly a *post-hoc reinterpretation* adopted after that failure, and section 5 gives the full rigor check—why the gap count is realization-, hyperparameter-, and key-fragile, and why we read the degradation as a direction rather than a majority-count finding. The uncertainty-taxonomy Mondrian add-on is a further negative control, not a co-equal finding (section 3): it is not a free win, and its calibration-size explanation is itself fragile—split-fragile (significant on only 4 of 20 split re-draws) and key-dependent (it clears the preregistered bar with the key as a feature but not without it)—an instance of the same sensitivity we document for the headline count, now applied reflexively to a candidate mechanistic explanation, which does not survive it.

Related work. A 2026 wave of work independently corroborates that tabular-FM accuracy does not buy reliability. De Melo Costa et al. (2026) report that tabular FMs beat GBDTs on AUC yet have worse conditional conformal coverage on 112 TALENT datasets—the same “high performance, low reliability” headline—but do not condition on split protocol (random vs. grouped/temporal), report risk–coverage/AURC, or test an abstention repair. The BeyondArena benchmark (Purucker et al., 2026) evaluates tabular FMs across IID/temporal/grouped splits on 142 datasets and finds FMs lead only on small IID data, but reports accuracy/ranking only—no calibration, conformal coverage, or selective prediction. Cheng et al. (2025) show TabPFN v2 degrades under distribution shift, and Helli et al. (2024) propose SCM priors to improve its temporal-shift calibration—the “FM auto-repairs shift” counter-hypothesis that our grouped-split arms refute. Two further 2026 entries study TabPFN uncertainty through decomposition (Fortini et al., 2026) and proper scoring rules (Landsgesell et al., 2026), but neither conditions on split protocol. A parallel line builds tabular out-of-distribution generalization benchmarks—TableShift (Gardner et al., 2023), WhyShift (Liu et al., 2023), and Wild-Tab (Kolesnikov, 2023)—that curate real distribution shifts and measure accuracy/robustness; we complement them by conditioning the same split types on reliability (risk–coverage and conformal coverage) and adding an abstention analysis, rather than reporting predictive performance alone. A closely related 2025–26 line develops conformal predictors that restore coverage under (sub)population shift—conformal prediction adaptive to unknown subpopulation shifts (Wang et al., 2025) and audited conformal prediction under unknown distribution shift (Zhou et al., 2026); these are exactly the coverage-side remedy for the exchangeability break we diagnose (section 2); we now evaluate the two standard distribution-free remedies (weighted / covariate-shift conformal and inferred-group Mondrian) in an exploratory coverage-remedy audit (section 3), finding that neither robustly restores group-conditional coverage under leave-group-out, while these stronger subpopulation- and audited-CP estimators remain untested here and an open follow-up rather than a competitor—our object is the split-conditioned reliability diagnostic plus a selective-abstention analysis, not a new coverage-restoring estimator. Two threads in this journal’s own recent literature bracket that object: Campagner et al. (2021) establish the isomorphisms between conformal prediction and three-way (cautious) decision-making, and He et al. (2025) couple conformal and

selective classification through a conjunction-subspaces test—placing the conformal-plus-abstention pairing we audit squarely in scope—while reject-option methods for tabular decision pipelines, such as credit-scoring reject inference (Shen et al., 2022), relevance-aggregation interpretability for tabular networks (Grisci et al., 2021), and streaming conformal prediction for evolving data (Tanha et al., 2022), map the surrounding tabular-reliability terrain. On the coverage side, adaptive conformal inference retunes the miscoverage level online to track a drifting distribution (Gibbs and Candès, 2021), and the reach of distribution-free guarantees toward the group-conditional coverage our leave-group-out split voids is delimited by conditional-validity analyses of inductive conformal predictors (Vovk, 2013) and by the jackknife+ (Barber et al., 2021). For the tabular substrate itself, controlled comparisons keep finding gradient-boosted trees competitive with or ahead of deep networks on typical tabular data (Grinsztajn et al., 2022; Gorishniy et al., 2021)—motivating our GBM-anchored roster—and convex-space oversampling has been proposed for the small, imbalanced tables where calibration data are scarcest (Schultz et al., 2024). Recent surveys of uncertainty quantification in deep models (Abdar et al., 2021; Gawlikowski et al., 2023) situate selective reliability within this broader reliability agenda. Against this backdrop, our contribution is a single, precisely scoped one, and we state its defense against a near-tautological reading up front rather than leave it implicit: we package the split-type $\times$ reliability $\times$ selective-abstention object with preregistered nulls and evaluate it across gradient-boosted models and two tabular foundation models, showing that one uncertainty-abstention recipe substantially improves selective reliability model-agnostically. Beating random deferral alone would be nearly definitional—any weakly informative score clears that bar—so the load-bearing claims are that the advantage holds under grouped and temporal shift, that no tabular foundation model auto-repairs it, and that the free score matches a budget-matched tuned selector at no extra cost, on both the exploratory suite and a preregistered held-out confirmatory replication (section 3), while we honestly scope which counts are realization-robust. We do not claim the four configurations are four independent confirmations—the two GBMs are strongly correlated (section 3), so they count as one architecture family plus two FM architectures. Methodologically, this is an empirical reliability/reproducibility study: we root-cause a cross-run irreproducibility to a shared-RNG advancement order (section 5) and replace single-split counts with repeated-draw percentile CIs, treating that discipline as a first-class result rather than an appendix. The conformal machinery we audit is standard split conformal (Vovk et al., 2005; Angelopoulos and Bates, 2023); abstention is selective classification (Geifman and El-Yaniv, 2017), which interacts with subgroup structure (Jones et al., 2021). Figure 1 summarizes the audit design used throughout: the frozen predictor and split-conformal wrapper are held fixed while only the split protocol and the coverage remedy vary, with the exploratory remedy audit quarantined from the confirmatory TableShift arm.




Figure 1: Overview of the reliability-collapse audit design: a keyed tabular dataset (entity or time key) is split randomly or by group/time (leave-group-out or rolling-origin); a frozen predictor (GBM or tabular FM, interchangeable) is wrapped in split-conformal prediction; a remedy ladder (weighted CP, Mondrian variants, tuned selector) is evaluated by worst-group coverage, yielding one of two verdicts: Restore (coverage returns to the 0.90 target, rare) or Fail (coverage remains below target, typical). The remedy audit is exploratory ($K=10$ split-repeats), quarantined from the confirmatory $K=20$ TableShift arm (no data overlap).

2 Data & Methods

Data. 14 OpenML datasets with a defensible entity key (e.g. `adult/native-country`, `nyc-taxi-green/location`, `sf-police-incidents/district`) or time key (e.g. `electricity`, `Airlines_DepDelay_10M`); selected by a 3-pass census (name $\rightarrow$ true-cardinality $\rightarrow$ hand-vetted) using the OpenML API (Vanschoren et al., 2014). This suite is deliberately keyed by construction rather than broad: the object of the audit is the reliability gap under a defensible deployment shift, which requires each dataset to carry a genuine entity/time key along which a leave-group-out / rolling-origin split is a real distributional change—most tabular benchmark datasets have none, and an arbitrary-column split is not a shift, so breadth and key-defensibility trade off directly. The 3-pass census reduced 145 candidates to 23 by name, 20 by true cardinality, and 14 hand-vetted with a reviewer-defensible key (frozen in a triage record); breadth is supplied separately by the 30-dataset arbitrary-key extension (section 3), kept out of the headline for exactly that reason.

Models. XGBoost (Chen and Guestrin, 2016) and LightGBM (Ke et al., 2017) with quantile-interval uncertainty (regression, whose conformalized form is the basis of conformalized quantile regression (Romano et al., 2019)) / $1 - \max p$ (classification), evaluated both at default hyperpa-

rameters and under a leakage-safe 12-trial randomized hyperparameter search (tuned only on the random-split training fold; the grouped test and calibration folds are never seen during tuning).

Foundation-model arms. On all 14 keyed datasets we additionally run two open-weights, token-free tabular foundation models: TabICLv2 (Qu et al., 2026) (`tabicl` package, `n_estimators=16`, GPU) and TabDPT (Ma et al., 2024) (`tabdpt` package, 8 ensemble members, GPU), each under the same split-conformal + abstention protocol; the same uncertainty score ($1 - \max p$ for classification; quantile width for regression) drives the abstention ranking. The TabICLv2 arm is run with the group key excluded and included (a leakage ablation) and against a de-staled Stage-1 GBM reference; a supplementary in-context scaling sweep (ensemble size $\in \{1, 4, 8, 16\} \times$ context length $\in \{200, 500, 1000, 2000\}$ on four datasets) and a 30-dataset breadth extension (the 14 keyed datasets plus 16 additional OpenML datasets grouped on an arbitrary categorical feature) are reported in section 3 as robustness checks, not as the headline.

Protocol. For each dataset and model we compute risk-coverage / AURC and split-conformal coverage ($\alpha = 0.10$; target 0.90) under a RANDOM split vs. a GROUPED (leave-group-out) / rolling-origin split, with 400-bootstrap CIs (Vovk et al., 2005; Angelopoulos and Bates, 2023). We then test an uncertainty-taxonomy Mondrian (Vovk et al., 2003) conformal predictor—one conformal quantile per decile of the model’s calibration uncertainty—against plain split-conformal. We use this confidence-conditional taxonomy rather than a group-keyed one because, under leave-group-out, the grouping key is absent from the calibration fold, so a group-keyed Mondrian provably degenerates to split conformal (we verified 0/28 deviation empirically). We then run a calibration-size diagnostic relating per-stratum calibration fold size to the uncertainty-taxonomy Mondrian’s coverage improvement. To characterize how sensitive the headline counts are to the particular split realization, we additionally repeat the GBM audit over $K=20\text{--}30$ independent grouped/rolling-origin split draws and report percentile CIs on the grouped-gap and repair counts (predeclared before the runs). Figures are generated from the result records.

Validity of the guarantee. Split-conformal coverage $\Pr(Y_{n+1} \in \hat{C}(X_{n+1})) \geq 1 - \alpha$ holds under a single assumption: the calibration points and the test point are *exchangeable* (Vovk et al., 2005; Lei et al., 2018). A random split satisfies this; a leave-group-out / rolling-origin split does not—the calibration fold is drawn from the training groups and the test point from disjoint held-out groups, so the two obey different group-conditional laws and the nonconformity scores are no longer exchangeable across that boundary. The finite-sample guarantee is then void, and coverage can fall below $1 - \alpha$ by the amount the score distribution shifts between groups—the covariate-shift / beyond-exchangeability regime of Tibshirani et al. (2019); Barber et al. (2023), whose recovery of coverage requires a known or estimated inter-group likelihood ratio that plain split conformal does not supply. The observed collapse (section 3) is therefore the predicted consequence of applying an exchangeability-dependent procedure across a group boundary that breaks exchangeability; the uncertainty-ranked abstention repair makes no exchangeability claim and restores low conditional risk on the retained set without it.

Reproducibility note. All XGBoost/LightGBM numbers in section 3 are read from two reconciled runs—a Stage-1 XGBoost-only audit and a Stage-2 combined two-model audit—regenerated under a per-context deterministic RNG discipline (an independent random-number stream seeded per dataset $\times$ purpose $\times$ model for splits, calibration carve-outs, and bootstraps; see section 5 for

the divergence this fixed). Under this discipline the two runs’ XGBoost arms produce identical per-dataset statistics (verified field-by-field on all 14 datasets), both GBM arms share identical splits and calibration folds, and reruns are deterministic. Result records carry provenance stamps (git SHA, seeds, package versions). Per-dataset one-sided bootstrap p -values are corrected across the 14 datasets with Holm–Bonferroni (FWER) and Benjamini–Hochberg (FDR) at $\alpha = 0.05$, using our reliability-analysis toolkit’s multiplicity-correction utilities.

3 Results

We flag each result as *confirmatory* (it meets a preregistered per-arm bar) or exploratory/ descriptive (a statistic we report without asserting it as a robust claim). The abstention repair meets its preregistered per-arm bar robustly throughout, including on the held-out TableShift replication below (H2 CONFIRM), and is confirmatory in every sense we use the word; the grouped-gap count—the only preregistered bar on the degradation side—is underpowered and is accordingly reported as a descriptive distribution rather than a confirmatory count. Section 5 gives the dedicated rigor check behind that labeling choice: our conjunctive Stage-2 gate on the gap count evaluated to KILL under the reconciled protocol, so the confirmatory/exploratory labels used throughout this section are a post-hoc reinterpretation, not a claim of overall preregistered success. Our contribution on the degradation axis is therefore the split-realization robustness analysis itself, not the gate verdict; we report the count below as a distribution over split draws and read the collapse direction as a descriptive observation, not a separately preregistered bar.

Grouped splits collapse reliability (descriptive); abstention substantially improves it (repair: per-arm confirmatory). Moving random→grouped degrades AURC suite-wide (e.g. *electricity* AURC $0.031 \rightarrow 0.224$; base error $0.117 \rightarrow 0.366$), while split-conformal under-coverage is concentrated on temporal splits (4/6 temporal model×dataset pairs undercover—*electricity*, *Airlines_DepDelay_10M*—vs. 4/22 grouped pairs), so the suite-wide collapse signal is the AURC one. Uncertainty-ranked abstention beats random deferral on 13 of 14 datasets (Fig. 2, 3); descriptively, it cuts integrated selective risk by a median of 33% (XGBoost) / 34% (LightGBM) relative to random deferral—a substantial improvement, though (lacking persisted risk–coverage curves at matched coverage) we do not quantify what fraction of the random→grouped gap this closes.




Figure 2: Per-dataset risk-coverage AURC under random vs. grouped/time split, shown as paired points on a log-scaled AURC axis (per-dataset AURC spans several orders of magnitude across classification and regression tasks, so a shared linear axis renders most datasets illegible): grouped splits degrade reliability.




Figure 3: Uncertainty-ranked abstention beats random deferral under the grouped split. Each point is one dataset’s repair ratio for one GBM (XGBoost, LightGBM) — the share of the random-deferral selective-risk AURC that abstention removes — with datasets ordered by ratio; points right of the dashed zero line improve on random deferral, and only a single model–dataset pair (Moneyball, XGBoost) falls at or below zero.

Model-agnostic across GBMs, with the repair robust and the count fragile. On a single reconciled split the grouped reliability gap excludes zero on 8/14 (XGBoost) and 7/14 (LightGBM) datasets with the group key retained as a feature, or on the more conservative, less-confounded key-excluded design—which we argue in section 5 is the less inflation-prone of the two—on only 6/14 and 5/14, both below the preregistered majority threshold of 8; repair-beats-random holds 13/14 for both GBMs in both key conditions (Fig. 4A).




Figure 4: Model-agnosticism across four configurations spanning three architecture families. A Grouped AUC for XGBoost vs. LightGBM across all 14 datasets (points near the diagonal confirm that the two GBMs behave near-identically, i.e. are effectively one GBDT family). B Split-conformal grouped coverage for all four configurations (XGBoost, LightGBM, and the TabICLv2 and TabDPT foundation models) on all 14 datasets; dashed line at target 0.90. Both FMs show the same collapse as the GBMs under grouped/time shift, with no auto-repair.

We report both key conditions side by side here rather than headlining the higher (key-included) count, and note that under the less-confounded design the gap-count component is a minority result—which only reinforces our thesis that the gap count is not the load-bearing claim. But that single-draw count is an *exploratory/descriptive* statistic, not the finding: over $K=20\text{--}30$ independent grouped split draws the grouped-gap count is a distribution with median 9/14 and 95% percentile interval $[6, 12]$ (XGBoost, key-included; LightGBM median 9, interval $[5, 11]$),¹ whose lower tail sits below the majority threshold—the gap count is majority-robust in only 75–83% of these split re-draws (which are independent re-partitions of the fixed 14-dataset panel, so the CI quantifies

¹Percentile-CI bounds are reported rounded to integer dataset counts (e.g. the XGBoost lower bound 6.7 is shown as 6); exact fractional bounds are in the split-repeats aggregate record.

split-realization variance, not sampling over a dataset population; the 83% upper end is the Stage-1 XGBoost arm, the 75% lower end the Stage-2 LightGBM arm), so “a majority of datasets show a gap” is not a realization-robust claim (Fig. 5).




Figure 5: Split-realization fragility of the four headline counts (LightGBM, key-included, K repeated grouped-split draws). Three of the four counts sit at or below the preregistered majority threshold of 8 (dashed line): the grouped-gap count (out of 14) is a broad distribution straddling the threshold—for LightGBM (shown) it is majority-robust in only 75% of draws (Stage-2 XGBoost 80%; combined range 75–83%)—the AURC-degrade count likewise straddles it, and the under-coverage count falls below it; only the repair-beats-random count sits at the top and never drops below the threshold (100% of draws). This is why we treat the repair, not the degradation-side counts, as the load-bearing finding.

The repair count, by contrast, is robust: its median is 13/14 and it exceeds the majority threshold in 100% of draws in every arm, with a 95% interval of [12, 14] for the Stage-1 XGBoost arm ($K=30$, the same arm as the 83% gap figure above) and [13, 14]/[12, 14] for the Stage-2 XGBoost/LightGBM arms ($K=20$). Hyperparameter tuning moves the single-draw count without a consistent sign (XGBoost $8 \rightarrow 6/14$, LightGBM $7 \rightarrow 8/14$; a wash on the majority question) while leaving repair at 13/14 (XGBoost) and 12/14 (LightGBM)—so the gap count is neither realization- nor hyperparameter-robust, whereas the repair is both. Multiple-comparison correction across the 14 datasets leaves the repair claim intact—13/14 datasets survive Holm–Bonferroni for both GBMs—while the AURC-degradation tests survive Holm on 6/14 (XGBoost) and 5/14 (LightGBM) and BH-FDR on 7/14 and 6/14. (A key-excluded robustness ablation—dropping the leave-group-out/rolling-origin key from the model features—lowers the single-draw gap counts to 6/14 XGBoost and 5/14 LightGBM while repair-beats-random stays at 13/14 for both; see section 5.)

Uncertainty-taxonomy Mondrian does not help (negative control). The uncertainty-taxonomy Mondrian add-on improves grouped coverage on 12/28 model×dataset pairs by point estimate but only 7/28 with CI excluding zero (median pooled improvement -0.002), well below the preregistered majority of 15 pairs, so it fails the Stage-2 Mondrian criterion. Its candidate mechanism, per-bin calibration fold size, is itself a case of the paper’s split-sensitivity lesson: on a single frozen split under our reconciled per-context RNG discipline, the clean-threshold criterion (balanced accuracy ≥ 0.75 , Mann–Whitney $p < 0.05$) passes in the key-included condition (bacc = 0.854, $p = 0.0011$, threshold 352 calibration observations; median per-bin size $\rho = 0.497$ [CI 0.10, 0.83], minimum per-bin size $\rho = 0.518$, $n_{\text{cal}} \rho = 0.461$, all Spearman CIs excluding zero),

but does not survive split-repeat robustness (median bacc = 0.64 [CI 0.55, 0.83] across 20 re-draws, significant on only 4/20, median $p = 0.21$ —the frozen 0.854/0.0011 is the favourable tail, not the typical case) and collapses entirely without the group/time key as a feature (zero predictors with CIs excluding zero, best separation bacc = 0.548, $p = 0.867$, chance level, preregistered verdict KILL; Fig. 6). We report the mechanism as real but split-fragile and key-dependent, not a confirmed explanation, per the paper’s own split-sensitivity standard.




Figure 6: Why Mondrian conformal fails, and why that explanation is key-dependent (reconciled RNG). A Key-included: per-bin calibration fold size predicts Mondrian’s coverage improvement (Spearman $\rho = 0.50$, CI excluding zero; clean-threshold separation bacc = 0.854, $p = 0.0011$)—the mechanism passes the preregistered bar. B Key-excluded: dropping the group key from the model features collapses the relationship ($\rho = -0.12$, no significant predictors, threshold bacc = 0.548), flipping the preregistered verdict to KILL. The calibration-size mechanism is real but key-dependent—the paper’s protocol-sensitivity finding applied to its own mechanism.

Is the repair a low bar? A learned selector does not beat the free uncertainty score

To check that abstention’s advantage over random deferral is not merely the low bar of beating a coin flip, we compare the free model-native uncertainty score against a learned post-hoc selector—a small untuned (default-hyperparameter, no search) gradient-boosted regressor trained on the calibration fold to predict per-example loss from prediction-level features (max-probability, entropy, margin for classification; interval width and endpoints for regression), using its predicted risk as the abstention score—in the identical risk–coverage frame (same test folds, calibration carve-out, and

AURC functions; the selector never sees the test fold). Table 1 summarizes the head-to-head over the 14 keyed datasets under the grouped split (XGBoost). The free uncertainty score is not a low bar: it reproduces the repair (beats random deferral on 13/14) and matches the trained selector, with statistically indistinguishable median repair ratios, both far below the oracle ceiling (Table 1; oracle ranks by the true loss). Throughout, the *repair ratio* is the fraction of the random-deferral selective-risk AURC that uncertainty-ranked abstention removes, $(\text{AURC}_{\text{random}} - \text{AURC}_{\text{abstain}}) / \text{AURC}_{\text{random}}$: higher is better, and zero means no improvement over random deferral. In the paired head-to-head the uncertainty score wins on 8/14 datasets and is never beaten on classification; the learned selector wins on only 3/14, all regression (**Airlines_DepDelay_10M**, **nyc-taxi-green-dec-2016**, **house_sales**), where it can exploit interval geometry beyond raw width, with 3 ties. The learned selector also needs calibration errors to fit, which dominant-group splits barely supply (**adult** and **cylinder-bands** leave only 56 and 45 calibration rows), a structural disadvantage the free score does not share. We scope this claim honestly: the selector is untuned and gets none of the 12-trial hyperparameter budget our GBMs receive, so the head-to-head is asymmetric in the free score’s favour, and the pattern is better read as “the free score wins where calibration errors are scarce (classification, small folds) and the selector wins where they are abundant (the three data-rich regression datasets)” than as a global tie—the equal medians average over those two regimes. We tested the obvious objection directly: giving the learned selector the same 12-trial randomized-search budget the GBMs receive (a budget-matched, exploratory head-to-head over $K=10$ grouped split-repeats with $N=400$ paired bootstraps) does not overturn the verdict—the free uncertainty score still ties-or-beats the budget-matched tuned selector on an aggregate median of 11.5 of the 14 datasets (min 9, max 13 across the ten split-repeats), and at the anchored single draw (k_0) on 9 of the 12 non-degenerate datasets (the remaining 2—**adult** and **cylinder-bands**—excluded because the tuned selector is untrainable on their dominant-group calibration fold). Our claim is therefore the weaker, defensible one—the free uncertainty score is not a low bar (it is never beaten on classification, ties on median, and is not overturned by a budget-matched tuned selector), not that no learned selector can ever win. This arm is XGBoost-only; LightGBM behaves near-identically throughout, so the verdict transfers.

Table 1: A learned post-hoc abstention selector vs. the free uncertainty score vs. random deferral (grouped split, 14 keyed datasets, XGBoost). Head-to-head is a paired test-set bootstrap on the repair-ratio difference; “wins” means the difference CI excludes zero. Numbers from the learned-selector baseline run.

Comparison / statistic	value
uncertainty abstention beats random deferral	13/14
learned selector beats random deferral	12/14
head-to-head: uncertainty wins	8/14
head-to-head: learned selector wins	3/14 (all regression)
head-to-head: tie	3/14
median repair ratio — uncertainty	0.339
median repair ratio — learned selector	0.338
median repair ratio — oracle (true-loss ranking)	0.745

Foundation-model arms: TabICLv2 and TabDPT collapse the same way

To pre-empt the question of whether a tabular foundation model auto-corrects the reliability gap, we ran two open-weights, token-free tabular FMs—TabICLv2 (Qu et al., 2026) (`tabicl`,

`n_estimators=16`) and TabDPT (Ma et al., 2024) (`tabdpt`, 8 ensemble members)—under the identical split-conformal + abstention protocol on all 14 keyed OpenML datasets, each with a per-context reconciled RNG and a de-staled Stage-1 GBM reference. Table 2 reports the TabICLv2 (leakage-ablated) numbers; Fig. 4B shows the four-model coverage comparison.

Table 2: TabICLv2 FM arm, leakage-ablated (`n_estimators=16`, key retained), all 14 datasets: split-conformal coverage ($\alpha = 0.10$, target 0.90) and grouped repair ratio under random vs. grouped/time split. Numbers from the leakage-ablated TabICLv2 arm; 400-bootstrap CIs. Undercover= grouped coverage < 0.88 (target -0.02 slack). The gap column is the coverage-gap point estimate (cov R – cov G computed before rounding), so it can differ from the displayed three-decimal subtraction by one unit in the last place.

Dataset	key	kind	cov R	cov G	gap	repair ratio (grouped)
electricity	date	time	0.898	0.634	+0.264	0.417 [0.396, 0.439]
eucalyptus	Year	time	0.904	0.838	+0.067	0.503 [0.400, 0.588]
adult	native-country	group	0.902	0.858	+0.043	0.703 [0.696, 0.710]
cylinder-bands	paper_mill_location	group	0.876	0.765	+0.111	0.337 [0.219, 0.451]
churn	state	group	0.907	0.913	−0.006	0.443 [0.247, 0.593]
Moneyball	Team	group	0.905	0.889	+0.016	0.064 [−0.011, 0.130]
kick	VNZIP1	group	0.895	0.893	+0.002	0.489 [0.447, 0.525]
black_friday	City_Category	group	0.905	0.899	+0.007	0.295 [0.287, 0.303]
house_prices_nominal	Neighborhood	group	0.952	0.940	+0.013	0.264 [0.204, 0.313]
colleges	state	group	0.914	0.891	+0.023	0.400 [0.381, 0.421]
Airlines_DepDelay	DayofMonth	time	0.906	0.234	+0.672	0.195 [0.184, 0.204]
nyc-taxi-green	PULocationID	group	0.901	0.895	+0.006	0.378 [0.358, 0.400]
house_sales	zipcode	group	0.909	0.817	+0.092	0.440 [0.422, 0.458]
sf-police-incidents	PdDistrict	group	0.894	0.891	+0.003	0.241 [0.194, 0.285]

The FM collapses under group/time shift across the full 14-dataset suite: for TabICLv2, AURC degrades (CI excluding 0) on 6/14 datasets, grouped conformal coverage falls below the literal 0.90 target on 12/14 datasets (median grouped coverage 0.890; on the stricter count that requires missing the target by more than the 0.02 slack used in Table 2, 6/14), and abstention-repair beats random deferral on 13/14 (median grouped repair ratio 0.389). Dropping the group key from the FM’s features (leakage ablation) leaves the repair at 13/14 and lowers the AURC-degrade count to 4/14, consistent with the GBM key ablation.

A second FM replicates it. TabDPT shows the same pattern: repair beats random on 13/14 datasets (both key conditions), grouped AURC degrades on 5/14, and median grouped coverage is 0.890—so the collapse-plus-repair story now reproduces across four configurations spanning three architecture families (GBDT plus two distinct tabular-FM architectures; Fig. 4B), no FM auto-repairing the gap. The hardest shifts remain temporal: `Airlines_DepDelay_10M` grouped coverage collapses to 0.234 and `electricity` to 0.634 (TabICLv2), comparable to or worse than the reconciled XGBoost baseline, confirming that the FMs require the same conformal + abstention recipe—not an auto-repair. These leakage-ablated `n_estimators=16` numbers, with a de-staled reference, supersede an earlier `n_estimators=4` pass whose 7/14 regression datasets had failed to download the FM regressor checkpoint in-container; the two passes agree qualitatively on the shared classification datasets (see section 5).

Robustness checks. A 30-dataset breadth extension (the 14 keyed datasets plus 16 OpenML datasets grouped on an arbitrary categorical feature, so the “grouped shift” is less meaningful there) keeps abstention-repair ahead of random on 25/29 completed datasets for TabICLv2 (the single non-completing dataset is `PhishingWebsites`, dropped at the arbitrary-key assignment gate before any model ran), and an in-context scaling sweep (ensemble $\times$ context length, 64 cells) shows AURC improving smoothly with more context without altering the grouped-collapse conclusion; we report both as appendix-level robustness, not as the keyed-suite headline.

Generalization to a TableShift-class benchmark (Folktables/ACS)

We test whether the model-agnostic finding survives on a canonical distribution-shift suite. The `tableshift` package ships no pip wheel, but it installs from source in an isolated environment and its 11 open tasks are credential-free (verified end-to-end by a loader smoke on `diabetes_readmission`). Two of these 11 credential-free tasks, `assistsments` and `college_scorecard`, proved Kaggle-API-gated at data-staging time and were excluded before any confirmatory number was computed, leaving 9 open tasks that enter the confirmatory run reported below, which resolves the held-out `ood_test` fold’s often-single-domain endpoint ambiguity with a frozen evaluability rule fixed ahead of time. For the present, complementary generalization arm we use the open TableShift-class substitute (Gardner et al., 2023): four Folktables/ACS (Ding et al., 2021) binary tasks (`ACIncome`, `ACSPublicCoverage`, `ACSMobility`, `ACSEmployment`) over US Census ACS PUMS 2018, with a leave-states-out spatial OOD split (held-out {MA, GA, MI}) against a random in-domain split, same protocol as above (XGBoost + LightGBM, split-conformal $\alpha = 0.10$, 400-bootstrap CIs).

Abstention repair replicates. Uncertainty-ranked abstention beats random deferral under OOD on 4/4 tasks for both GBMs, every CI excluding 0 (median OOD repair ratio 0.539 XGBoost, 0.579 LightGBM; Fig. 7A).

The OOD reliability gap manifests as a subgroup coverage gap (exploratory generalization). We flag this arm as exploratory: the primary, preregistered marginal criterion (AURC-degrade or marginal under-coverage) replicated on only 1/4 tasks per GBM under this milder spatial shift, so the finding below rests on a post-hoc subgroup criterion and should be read as hypothesis-generating, not confirmatory. On this milder ACS spatial shift the marginal-AURC degradation is weak (only 1/4 tasks per GBM degrade with CI > 0; median AURC Δ 0.003 XGBoost, -0.002 LightGBM); instead the gap appears as a RAC1P race-subgroup conformal coverage gap—worst-minus-best is negative on all 8 (model, task) cells, with the gap CI excluding 0 on 4/4 tasks for both GBMs (median -0.033 XGBoost, -0.027 LightGBM; Fig. 7B)—but this 4/4 is the post-hoc, broadened subgroup criterion, not the preregistered marginal one. The honest reading is that only the repair generalizes to this milder ACS shift (the abstention improvement replicates 4/4); the marginal reliability collapse does not (AURC-degrade / under-coverage on just 1/4 tasks per GBM), and the RAC1P subgroup-coverage-gap 4/4 survives only because the success criterion was broadened after the marginal one missed—so the collapse itself is benchmark-dependent and does not generalize here, only the repair does.

Note on the reliability-gap criterion. The OOD reliability-gap criterion here is a 3-way disjunction (AURC-degrade CI > 0 OR marginal under-coverage OR RAC1P subgroup-coverage-gap CI < 0), broadened from the primary 14-dataset study’s 2-way disjunction (AURC-degrade OR coverage-gap only), because the narrow criterion alone replicated on only 1/4 tasks per GBM under this milder spatial shift.




Figure 7: Generalization to a TableShift-class benchmark (Folktables/ACS, leave-states-out spatial OOD). A OOD abstention repair ratio per ACS task (XGBoost vs. LightGBM, 400-bootstrap CIs); the repair ratio is the fractional reduction in integrated selective risk vs. random deferral, so positive ratios (CI excluding 0) mean uncertainty abstention beats random deferral, replicating on 4/4 tasks for both GBMs. B RAC1P race-subgroup conformal coverage gap (worst–best); the OOD reliability gap manifests as a negative subgroup coverage gap (CI excluding 0 on 4/4 for both GBMs) rather than a marginal-AURC blow-up.

Exploratory audit: do distribution-free remedies restore group-conditional coverage? (No.)

We close the diagnostic by asking, exploratorily, whether the coverage side of the collapse can be repaired by a *distribution-free* conformal remedy that never sees the (absent) grouping key. We label this arm exploratory throughout, per its own preregistration: it re-uses the 14 keyed OpenML datasets that generated the hypothesis, so the confirmatory test is reserved for a held-out TableShift family the exploratory data cannot contaminate, and nothing here is read as confirmatory. It is run at $K=10$ grouped split-repeats—the preregistration specifies $K=20$ for the confirmatory version; we report $K=10$ here as the exploratory budget and flag the difference—with $N=400$ paired-bootstrap CIs and the GBM loss/uncertainty arrays bit-identical to the frozen backbone.

The disease. Under the grouped / rolling-origin split, plain split conformal—the paper’s headline recommendation—under-covers group-conditionally: median worst-group coverage 0.810 against the 0.90 target, median worst–best group gap 0.144. The marginal guarantee holds; the group-conditional one does not—the exchangeability break made concrete.

The remedy evaluation. We run the two standard distribution-free remedies—weighted / covariate-shift CP (Tibshirani et al., 2019) and inferred-group Mondrian (group-conditional CP on KMeans clusters of the covariate space, the split key dropped)—plus the preregistered KILL null (the uncertainty-decile Mondrian). A remedy restores coverage on a dataset only under a conjunctive bar: worst-group coverage ≥ 0.87 and a bootstrap-CI-positive group-gap reduction vs. split conformal and no set-size blow-up. The count of the 14 datasets where any remedy clears that bar has median 1 (min 0, max 3) across the ten split-repeats—never a majority (threshold 8). Per remedy, the median worst-group coverage at the anchored draw (k_0) is 0.831 (weighted), 0.817 (inferred-Mondrian) and 0.811 (unc-decile KILL null)—all below the 0.87 restoration threshold, so the negative already holds on the worst-group point estimate alone, before the conjunctive CI and set-size conditions are even applied.

The mechanism. This is not a failure of the remedies in general: weighted CP provably restores coverage under pure covariate shift, and inferred-Mondrian restores worst-group coverage under subpopulation heterogeneity—both are validated on synthetic data constructed so those assumptions hold exactly. Their failure here indicates that leave-group-out induces a shift beyond covariate-only (a label/concept component the covariate-space reweighting cannot capture), or from calibration-fold scarcity our design cannot separate, so we scope the negative narrowly to the two standard distribution-free remedies we implemented. The stronger subpopulation-shift and audited-CP estimators (Wang et al., 2025; Zhou et al., 2026) are not run here; we cite them as the open coverage-restoring question, not as implied-refuted. The honest prescription is therefore negative for coverage, positive for selection: no tested distribution-free remedy restores group-conditional coverage without the group key, while the free uncertainty score reliably ties-or-beats a budget-matched tuned selector (above).

Confirmatory coverage-remedy evaluation on the held-out TableShift roster

We now report the confirmatory test the exploratory remedy audit above was reserved for: a preregistered, held-out replication of H1 and H2 on TableShift’s own open task roster, disjoint from the 14 keyed OpenML datasets that generated the two hypotheses. Frozen ahead of any confirmatory number: **H1 (coverage restoration)**—some single remedy restores group-conditional coverage on a majority of evaluable held-out tasks, at the minimum across $K=20$ split-repeats; **H2 (selective prescription)**—the free uncertainty score ties or beats a budget-matched, 12-trial-tuned learned selector on a majority of evaluable tasks, robustly across K .

Roster and frozen endpoint rule. All 9 open (credential-free) tasks complete (9/9; $K=20$ split-repeats each, $N=400$ -bootstrap CIs). A single rule, frozen and applied mechanically before any confirmatory number was computed, assigns each task exactly one endpoint (Table 3): 3 tasks have too few domain groups for any group-conditional endpoint and are excluded-by-rule from the H1/H2 tally; the remaining 6 evaluable tasks split into 2 with a primary worst-group-domain endpoint and 4 with a fallback marginal-gap-plus-secondary-grouping endpoint. That split was chosen by the same frozen rule, and the reader should weigh it accordingly: 4 of the 6 evaluable tasks certify under the fallback (marginal-gap plus secondary-grouping) endpoint rather than the primary worst-group-domain endpoint that motivates the study, so the roster-level H1/H2 verdicts below are majority-carried by the fallback construct, not the primary one. The rule is endpoint-blind—it selects an endpoint from each task’s own group structure, not from which answer that endpoint would produce—so the fallback-heavy composition is a property of the roster, not a chosen-

to-flatter selection; read plainly, the confirmatory yield below is one confirmed negative on coverage (H1 REJECT) plus one confirmed tie on selection (H2 CONFIRM), not a positive method win.

Table 3: Frozen per-task endpoint decisions on the 9-task open, held-out TableShift roster (evaluability rule applied mechanically, before any confirmatory number was computed). The 3 excluded-by-rule tasks report only a descriptive marginal in-domain-vs-out-of-domain coverage gap (main text); the 6 evaluable tasks enter the H1/H2 tally in Table 4.

Task	Endpoint	n groups	In H1/H2 tally
diabetes_readmission	marginal gap only	0	excluded-by-rule
nhanes_lead	marginal gap only	0	excluded-by-rule
physionet	marginal gap only	0	excluded-by-rule
brfss_diabetes	primary worst-group domain	5	yes
brfss_blood_pressure	fallback (secondary: AGE5YR)	7	yes
acsincome	fallback (secondary: AGE)	71	yes
acspubcov	fallback (secondary: AGE)	50	yes
acsfoodstamps	fallback (secondary: AGE)	44	yes
acsunemployment	primary worst-group domain	15	yes

H1: no remedy robustly restores group-conditional coverage (REJECT). Under the leave-group-out split, plain split conformal under-covers group-conditionally on the evaluable roster (median worst-group coverage 0.797 against the 0.90 target, median worst–best group gap 0.135), replicating the exploratory audit’s diagnosis on a fully held-out benchmark. The count of evaluable tasks where any remedy clears the restores-coverage bar has minimum 3, median 3, maximum 4 across the 20 split-repeats—never reaching the majority threshold of 4 at the minimum, so H1 is REJECTed. Applying the same fragility scrutiny used elsewhere in this paper to the 14-dataset gap count, this REJECT itself rests on only 6 evaluable tasks (4 of 6 on the fallback endpoint, Table 3) and is not robust to which task is held out: recomputing the frozen $\lfloor n/2 \rfloor + 1$ majority rule at $n=5$ (leave-one-task-out) flips the verdict to CONFIRM if any of the three tasks where no remedy ever restores coverage across K (brfss_diabetes, brfss_blood_pressure, acsunemployment) is dropped, and leaves it REJECT if any of the three tasks where a remedy restores coverage at every repeat (acsincome, acspubcov, acsfoodstamps) is dropped. Per remedy (min/median/max across K): weighted CP 3/3.0/4, inferred-group Mondrian 0/0.0/0, uncertainty-decile Mondrian 0/0.0/1 (Table 4)—the same ordering as the exploratory audit, with weighted CP strongest and still short of a robust majority. A supplementary, non-decision-altering multiplicity correction (Holm–Bonferroni / Benjamini–Hochberg across the 6 evaluable tasks, anchored single-draw p -values) rejects the gap-reduction null on 5/6 tasks at $\alpha = 0.05$ (all but acsunemployment, $p = 0.165$): individual remedies do measurably shrink the gap on most tasks, just not far enough, and not consistently enough across K , to clear the conjunctive restoration bar on a majority.

H2: the free uncertainty score ties-or-beats the tuned selector on every task (CONFIRM). The uncertainty-ranked selector ties or beats the 12-trial-tuned learned selector on all 6/6 evaluable tasks, at every one of the $K=20$ split-repeats (minimum = median = maximum = 6)—a fully robust replication of the exploratory audit’s finding, with no exception across any task or repeat. The matching supplementary multiplicity correction on the tuned-minus-uncertainty repair-ratio delta finds $p = 1.000$ on all 6 tasks (the tuned selector never wins). H2 is CONFIRMed.

Table 4: Confirmatory H1/H2 verdicts on the 6 evaluable, held-out TableShift tasks ($K=20$ split-repeats, $N=400$ -bootstrap CIs). Min/median/max are computed across the 20 split-repeats; majority threshold= 4.

Comparison / statistic	value
median split-conformal worst-group coverage (target 0.90)	0.797
median split-conformal worst–best group gap	0.135
any-remedy restores coverage (min / median / max across K)	3 / 3.0 / 4
weighted CP	3 / 3.0 / 4
inferred-group Mondrian	0 / 0.0 / 0
uncertainty-decile Mondrian	0 / 0.0 / 1
H1 (coverage restoration)	REJECT
uncertainty ties-or-beats tuned selector (min / median / max across K)	6 / 6.0 / 6
H2 (selective prescription)	CONFIRM

Descriptive: the marginal id-vs-ood coverage gap where available. TableShift exposes an additional in-domain held-out (`id_test`) split for only 3 of the 9 tasks; where available, the marginal (non-group-conditional) coverage gap under plain split conformal is `physionet` +0.094, `brfss_diabetes` +0.040, and `diabetes_readmission` −0.018—two of three tasks show the same direction (out-of-domain coverage below in-domain) as the primary 14-dataset study’s headline collapse, with `diabetes_readmission`’s small negative gap the exception.

Mechanical loader adapters (disclosed deviations). Running the frozen endpoint rule and remedy ladder on TableShift’s own loaders required three narrow, mechanical adapters, none of which touch the frozen decision rule or the remedy code itself: (i) a categorical domain-column passthrough for the ACS tasks, reproducing TableShift’s own (empty) preprocessing configuration plus only the domain column; (ii) a case-corrected, isolated cache directory for the NHANES loader, working around an uppercase/lowercase filename mismatch in TableShift’s own downloader that would otherwise silently substitute a different survey wave lacking the needed fields; (iii) column-restricted reads for the wide ACS person/household files, reducing memory footprint without altering which columns are used. Each adapter was validated by reproducing an already-completed task and confirming identical results (every field matched except wall-clock timing) before being adopted on the remaining tasks; the frozen evaluability/endpoint rule itself was never modified.

4 Discussion

The reliable, simple recipe for the selective reliability of the retained set under grouped tabular shift is uncertainty abstention on top of plain split-conformal, and it is model-agnostic across four open, token-free configurations spanning three architecture families: two gradient-boosted models (which behave near-identically, effectively one GBDT family) and two distinct tabular-FM architectures. We are explicit that this recipe targets selective reliability, not marginal coverage: uncertainty abstention makes no coverage guarantee, and the split-conformal coverage that collapses under group shift (e.g. `Airlines_DepDelay_10M` to 0.234) is not restored by it. Restoring marginal coverage under this exchangeability break is the province of weighted or adaptive conformal (Tibshirani et al., 2019; Barber et al., 2023); neither of the two standard distribution-free remedies we evaluate (weighted / covariate-shift CP and inferred-group Mondrian) robustly restores it without the absent group key (section 3), and the stronger subpopulation- and audited-CP methods (Wang et al., 2025;

Zhou et al., 2026) remain an open, still-untested follow-up; accordingly we do not recommend bare split-conformal as a coverage fix. A preregistered confirmatory replication on the held-out TableShift roster reaches the same verdict (H1 REJECT; H2 CONFIRM without exception across any task or repeat; section 3)—so the negative-coverage/positive-selection prescription is not an artifact of the 14-dataset convenience suite. Neither FM (TabICLv2 (Qu et al., 2026), TabDPT (Ma et al., 2024)) auto-repairs the reliability gap: both show the same collapse under group/time shift as XGBoost and LightGBM, and require the same conformal + abstention recipe. What we deliberately do not claim is a stable “majority of datasets show a gap” count: over repeated split draws, hyperparameter tuning, and key inclusion, that count ranges over roughly 5–12 of 14 and is majority-robust in only three-quarters to four-fifths of draws with the group key included, falling to roughly half with the key excluded (section 5), whereas the repair beats random deferral on 13/14 in every arm and every draw. The uncertainty-taxonomy Mondrian add-on is not a free win: it needs large per-stratum calibration, which grouped-shift tabular data rarely provides. The calibration-size explanation for that failure (section 3, Fig. 6) clears our preregistered bar only on a single frozen split with the group key retained, does not survive our own split-repeat robustness check, and collapses to chance once the key is dropped; we therefore treat it as a split-fragile, key-dependent association rather than an established mechanism, and recommend reporting all conditions rather than the favorable one alone.

5 Limitations

This section is this paper’s dedicated rigor check: the preregistered Stage-2 gate’s KILL verdict, the split-realization fragility behind Fig. 5, and the key-inclusion sensitivity of the headline counts are stated here once, in full, rather than scattered across the paper, because the honest scope of a preregistered study belongs in one place a reader can check against the confirmatory claims above.

Scope and future work. The 14 keyed OpenML datasets are a convenience suite; we have already extended to the open TableShift-class benchmark (Folktables/ACS) above, where the abstention-repair finding replicates and the OOD reliability gap surfaces as a RAC1P subgroup coverage gap. As noted above (section 3), the ACS reliability-gap criterion was broadened to a 3-way disjunction because the narrow 2-way criterion alone replicated on only 1/4 tasks per GBM under this milder spatial shift. The FM roster now covers all 14 keyed datasets for both TabICLv2 (leakage-ablated, `n_estimators=16`, de-staled reference) and TabDPT, so the “larger-ensemble configurations remain future work” caveat of earlier drafts is discharged; an earlier `n_estimators=4` pass had left the 7 regression datasets incomplete because the FM regressor checkpoint failed to download in-container, which a checkpoint pre-fetch resolved. The FM roster still omits a token-gated TabPFN-v2/2.5 arm (Hollmann et al., 2025; Grinsztajn et al., 2025), parked behind a license review; we defend the current roster on license-cleanliness (all four configurations are open and token-free) and note that adding TabPFN is the cheapest response to a “where is TabPFN” referee question. The credentialed TableShift medical tasks (e.g. MIMIC/ICU loaders) remain future work behind their data licenses.

Cross-run reproducibility of XGBoost — root-caused and fixed; the cost is one honest count revision. The two independent XGBoost runs behind earlier drafts (the Stage-1 and Stage-2 audits) diverged sharply on some datasets (e.g. `adult` grouped repair-ratio 0.635 vs. 0.119; grouped-gap count 8/14 vs. 10/14) despite fully seeded, internally deterministic scripts. The cause was a shared-RNG advancement order: both scripts drove every stochastic step through a single module-level generator, so Stage-2’s second (LightGBM) arm advanced the stream before the next

dataset’s splits were drawn. Dataset #1 (**electricity**) matched bit-for-bit across the two runs while every later dataset diverged, confirming the two runs were evaluating different split realizations rather than exhibiting model irreproducibility. We replaced the shared stream with per-context deterministic RNGs (seeded per dataset $\times$ purpose $\times$ model): the two XGBoost arms now agree exactly on all 14 per-dataset records, and both GBMs share identical splits and calibration folds. Under the reconciled protocol the headline counts are those reported in section 3 (gap 8/14 XGBoost, 7/14 LightGBM; repair 13/14 both). LightGBM’s gap count now sits one dataset below the preregistered majority threshold, so the conjunctive Stage-2 gate (gap-count Q1 $\wedge$ Mondrian Q2) evaluates to KILL under the reconciled protocol even though the repair sub-criterion replicates strongly; the confirmatory/exploratory split used throughout is accordingly a post-hoc reinterpretation, not an overall preregistered success (section 3). The gate is underpowered on the gap arm: Stage-1 (XGBoost) passed it at exactly the majority threshold (8/14, a knife-edge), a single-dataset flip is enough to fail it, and single-draw counts across the pre- and post-fix realizations have in fact ranged 7–10 of 14—fig. 5 and section 3 give the full split-realization distribution behind that fragility. A verdict this sensitive to one re-draw carries little information, so we treat the Holm-robust repair as the confirmatory finding and the degradation direction as descriptive, not the gate outcome itself. The headline FM arm (**n_estimators=16** leakage ablation) carries a de-staled Stage-1 GBM reference; only the superseded **n_estimators=4** pass retains the pre-reconciliation reference and is no longer cited for any headline number.

Group-key feature leakage — bounded by a key-excluded ablation. In the main reported runs the leave-group-out/rolling-origin key (e.g. **native-country**, **date**) is retained as a label-encoded model feature. The Stage-2 preregistration described this choice as conservative for the degradation claim, on the reasoning that seeing the key would make grouped prediction easier; that reasoning understates the risk in the other direction. Because leave-group-out test keys take categorical values never seen during training, a model that has learned to rely on key-specific patterns in-sample can behave erratically out-of-sample on the encoded key alone, so retaining the key can inflate, not just dampen, the apparent grouped-split collapse—confounding genuine group/time shift with an ID-generalization artifact. We re-ran the Stage-1 and Stage-2 pipelines with the split key dropped from the feature encoder (the key-excluded runs); we also re-ran both FM arms under the key-excluded ablation (TabICLv2 and TabDPT, repair 13/14 in both conditions), while the TableShift/ACS arm is out of scope for this specific Table 1 headline (its own separate reliability-gap criterion is discussed above). Under the reconciled RNG protocol, excluding the key lowers the grouped-gap count from 8/14 to 6/14 (XGBoost; identical in the Stage-1 and Stage-2 runs, which now agree exactly) and from 7/14 to 5/14 (LightGBM)—below the preregistered majority threshold of 8 for both GBMs—while the repair-beats-random count is unchanged at 13/14 for both. We read this as evidence that key leakage inflates the grouped-gap count (roughly two datasets’ worth for each GBM), on top of the split-realization sensitivity documented above. The robust core claim is therefore the repair result—uncertainty abstention beats random deferral on 13/14 datasets in every arm (key-included, key-excluded, both GBMs, and both FMs; the GBM/FM correction asymmetry behind that shorthand is detailed below)—while the grouped-degradation majority count should be read as realization- and design-sensitive (range 5–8 of 14 across the reconciled key-included/-excluded runs) pending repeated-split-draw CIs. Section 3’s headline numbers cite the key-included reconciled runs; the key-excluded counts above are the conservative lower bound.

FM repair claims rest on CIs, not the Holm-corrected bar (limitation). The GBM repair claim clears Holm–Bonferroni across all 14 datasets; the two FM arms’ repair claim does not

undergo the same correction, because per-dataset p -values for a Holm pass on the FMs were not persisted, so the FM 13/14 is corroborated only by 400-bootstrap CIs excluding zero—a CI-level, not a FWER-level, result. We flag this asymmetry explicitly as a limitation: wherever the “13/14 in every arm” shorthand appears in this paper, the two claims—GBM = Holm-corrected, FM = CI-only—should be read as verbally distinct, never as an equally multiplicity-corrected result across all four configurations.

Data & Code Availability

All 14 datasets are public OpenML ([Vanschoren et al., 2014](#)) (no auth); models XGBoost ([Chen and Guestrin, 2016](#)), LightGBM ([Ke et al., 2017](#)), TabICLv2 ([Qu et al., 2026](#)) (`tabicl` package), and TabDPT ([Ma et al., 2024](#)) (`tabdpt` package), all open-weights and token-free. The frozen result records and the R figure scripts (regenerating every figure from those records) accompany this draft; reliability metrics use the audit implementation vendored with the released code. The development repository is <https://github.com/PeterPonyu/structured-data-fm-reliability-research>; the code is archived at Zenodo under DOI 10.5281/zenodo.21130297, which is currently reserved on a draft deposition and becomes active when the record is published.

References

- Moloud Abdar, Farhad Pourpanah, Sadiq Hussain, Dana Rezazadegan, Li Liu, Mohammad Ghavamzadeh, Paul Fieguth, Xiaochun Cao, Abbas Khosravi, U. Rajendra Acharya, Vladimir Makarenkov, and Saeid Nahavandi. A review of uncertainty quantification in deep learning: Techniques, applications and challenges. *Information Fusion*, 76:243–297, 2021. doi: 10.1016/j.inffus.2021.05.008.
- Anastasios N. Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023. doi: 10.1561/2200000101.
- Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. Predictive inference with the jackknife+. *The Annals of Statistics*, 49(1):486–507, 2021. doi: 10.1214/20-AOS1965.
- Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. Conformal prediction beyond exchangeability. *The Annals of Statistics*, 51(2):816–845, 2023. doi: 10.1214/23-AOS2276.
- Andrea Campagner, Federico Cabitza, Pedro Berjano, and Davide Ciucci. Three-way decision and conformal prediction: Isomorphisms, differences and theoretical properties of cautious learning approaches. *Information Sciences*, 579:347–367, 2021. doi: 10.1016/j.ins.2021.08.009.
- Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016. doi: 10.1145/2939672.2939785.
- Zi-Jian Cheng, Zi-Yi Jia, Zhi Zhou, Yu-Feng Li, and Lan-Zhe Guo. Realistic evaluation of TabPFN v2 in open environments. arXiv:2505.16226, 2025.
- C. K. Chow. On optimum recognition error and reject tradeoff. *IEEE Transactions on Information Theory*, 16(1):41–46, 1970. doi: 10.1109/TIT.1970.1054406.

- Corinna Cortes, Giulia DeSalvo, and Mehryar Mohri. Learning with rejection. In *Algorithmic Learning Theory (ALT)*, Lecture Notes in Computer Science, pages 67–82, 2016. doi: 10.1007/978-3-319-46379-7_5.
- José Lucas De Melo Costa, Fabrice Popineau, Arpad Rimmel, and Bich-Liên Doan. High performance, low reliability: Uncertainty benchmarking for tabular foundation models. arXiv:2605.28554, 2026. ESANN 2026.
- Frances Ding, Moritz Hardt, John Miller, and Ludwig Schmidt. Retiring Adult: New datasets for fair machine learning. In *Advances in Neural Information Processing Systems 34 (NeurIPS)*, pages 6478–6490, 2021. Folktables package.
- Sandra Fortini, Kenyon Ng, Sonia Petrone, Judith Rousseau, and Susan Wei. A principled framework for uncertainty decomposition in TabPFN. arXiv:2602.04596, 2026.
- Josh Gardner, Zoran Popović, and Ludwig Schmidt. Benchmarking distribution shift in tabular data with TableShift. In *Advances in Neural Information Processing Systems 36 (NeurIPS), Datasets and Benchmarks Track*, 2023.
- Jakob Gawlikowski, Cedrique Rovile Njietcheu Tassi, Mohsin Ali, Jongseok Lee, Matthias Humt, Jianxiang Feng, Anna Kruspe, Rudolph Triebel, Peter Jung, Ribana Roscher, Muhammad Shahzad, Wen Yang, Richard Bamler, and Xiao Xiang Zhu. A survey of uncertainty in deep neural networks. *Artificial Intelligence Review*, 56(Suppl 1):1513–1589, 2023. doi: 10.1007/s10462-023-10562-9.
- Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems 30 (NeurIPS)*, pages 4878–4887, 2017.
- Isaac Gibbs and Emmanuel J. Candès. Adaptive conformal inference under distribution shift. arXiv:2106.00170, 2021. Advances in Neural Information Processing Systems 34 (NeurIPS).
- Yury Gorishniy, Ivan Rubachev, Valentin Khrulkov, and Artem Babenko. Revisiting deep learning models for tabular data. arXiv:2106.11959, 2021. Advances in Neural Information Processing Systems 34 (NeurIPS).
- Léo Grinsztajn, Edouard Oyallon, and Gaël Varoquaux. Why do tree-based models still outperform deep learning on tabular data? arXiv:2207.08815, 2022. Advances in Neural Information Processing Systems 35 (NeurIPS), Datasets and Benchmarks Track.
- Léo Grinsztajn, Klemens Flöge, Oscar Key, Felix Birkel, Philipp Jund, Brendan Roof, Benjamin Jäger, Dominik Safaric, Simone Alessi, Adrian Hayler, Mihir Manium, Rosen Yu, Felix Jablonski, Shi Bin Hoo, Anurag Garg, Jake Robertson, Magnus Bühler, Vladyslav Moroshan, Lennart Puercker, Clara Cornu, Lilly Charlotte Wehrhahn, Alessandro Bonetto, Bernhard Schölkopf, Sauraj Gambhir, Noah Hollmann, and Frank Hutter. TabPFN-2.5: Advancing the state of the art in tabular foundation models. arXiv:2511.08667, 2025. Prior Labs.
- Bruno Iochins Grisci, Mathias J. Krause, and Marcio Dorn. Relevance aggregation for neural networks interpretability and knowledge discovery on tabular data. *Information Sciences*, 559: 111–129, 2021. doi: 10.1016/j.ins.2021.01.052.
- Zengyou He, Zerun Li, Junjie Dong, Xinying Liu, Mudi Jiang, and Lianyu Hu. Conjunction subspaces test for conformal and selective classification. *Information Sciences*, 707:122037, 2025. doi: 10.1016/j.ins.2025.122037.

- Kai Helli, David Schnurr, Noah Hollmann, Samuel Müller, and Frank Hutter. Drift-resilient TabPFN: In-context learning temporal distribution shifts on tabular data. arXiv:2411.10634, 2024. NeurIPS 2024.
- Radu Herbei and Marten H. Wegkamp. Classification with reject option. *Canadian Journal of Statistics*, 34(4):709–721, 2006. doi: 10.1002/cjs.5550340410.
- Noah Hollmann, Samuel Müller, Lennart Purucker, Arjun Krishnakumar, Max Körfer, Shi Bin Hoo, Robin Tibor Schirrmester, and Frank Hutter. Accurate predictions on small data with a tabular foundation model. *Nature*, 637:319–326, 2025. doi: 10.1038/s41586-024-08328-6. TabPFN v2.
- Erik Jones, Shiori Sagawa, Pang Wei Koh, Ananya Kumar, and Percy Liang. Selective classification can magnify disparities across groups. In *International Conference on Learning Representations (ICLR)*, 2021.
- Faisal Kamiran, Sameen Mansha, Asim Karim, and Xiangliang Zhang. Exploiting reject option in classification for social discrimination control. *Information Sciences*, 425:18–33, 2018. doi: 10.1016/j.ins.2017.09.064.
- Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems 30 (NeurIPS)*, pages 3146–3154, 2017.
- Sergey Kolesnikov. Wild-Tab: A benchmark for out-of-distribution generalization in tabular regression. arXiv:2312.01792, 2023.
- Jonas Landsgesell, Pascal Knoll, and Tizian Wenzel. Distributional regression with tabular foundation models: Evaluating probabilistic predictions via proper scoring rules. arXiv:2603.08206, 2026.
- Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523): 1094–1111, 2018. doi: 10.1080/01621459.2017.1307116.
- Jiashuo Liu, Tianyu Wang, Peng Cui, and Hongseok Namkoong. On the need for a language describing distribution shifts: Illustrations on tabular datasets. In *Advances in Neural Information Processing Systems 36 (NeurIPS), Datasets and Benchmarks Track*, 2023. Introduces the WhyShift benchmark; arXiv:2307.05284.
- Junwei Ma, Valentin Thomas, Rasa Hosseinzadeh, Alex Labach, Hamidreza Kamkari, Jesse C. Cresswell, Keyvan Golestan, Guangwei Yu, Anthony L. Caterini, and Maksims Volkovs. TabDPT: Scaling tabular foundation models on real data. arXiv:2410.18164, 2024. <https://github.com/layer6ai-labs/TabDPT>.
- Harris Papadopoulos, Kostas Proedrou, Volodya Vovk, and Alex Gammerman. Inductive confidence machines for regression. In *Machine Learning: ECML 2002*, Lecture Notes in Computer Science, pages 345–356, 2002. doi: 10.1007/3-540-36755-1_29.
- Lennart Purucker, Andrej Tschalzev, Nick Erickson, Gioia Blayer, David Holzmüller, Alan Arazi, Alexander Pfefferle, Mustafa Tajjar, Gaël Varoquaux, and Frank Hutter. Beyond IID: How general are tabular foundation models, really? arXiv:2606.30410, 2026. Introduces the BeyondArena benchmark.

- Jingang Qu, David Holzmüller, Gaël Varoquaux, and Marine Le Morvan. TabICLv2: A better, faster, scalable, and open tabular foundation model. arXiv:2602.11139, 2026. <https://github.com/soda-inria/tabicl>.
- Yaniv Romano, Evan Patterson, and Emmanuel J. Candès. Conformalized quantile regression. arXiv:1905.03222, 2019. *Advances in Neural Information Processing Systems 32 (NeurIPS)*.
- Kristian Schultz, Saptarshi Bej, Waldemar Hahn, Markus Wolfien, Prashant Srivastava, and Olaf Wolkenhauer. ConvGeN: A convex space learning approach for deep-generative oversampling and imbalanced classification of small tabular datasets. *Pattern Recognition*, 147:110138, 2024. doi: 10.1016/j.patcog.2023.110138.
- Feng Shen, Zhiyuan Yang, Xingchao Zhao, and Dao Lan. Reject inference in credit scoring using a three-way decision and safe semi-supervised support vector machine. *Information Sciences*, 606: 614–627, 2022. doi: 10.1016/j.ins.2022.05.067.
- Jafar Tanha, Negin Samadi, Yousef Abdi, and Nazila Razzaghi-Asl. CPSSDS: Conformal prediction for semi-supervised classification on data streams. *Information Sciences*, 584:212–234, 2022. doi: 10.1016/j.ins.2021.10.068.
- Ryan J. Tibshirani, Rina Foygel Barber, Emmanuel J. Candès, and Aaditya Ramdas. Conformal prediction under covariate shift. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019.
- Joaquin Vanschoren, Jan N. van Rijn, Bernd Bischl, and Luís Torgo. OpenML: Networked science in machine learning. *ACM SIGKDD Explorations Newsletter*, 15(2):49–60, 2014. doi: 10.1145/2641190.2641198.
- Vladimir Vovk. Conditional validity of inductive conformal predictors. *Machine Learning*, 92(2–3): 349–376, 2013. doi: 10.1007/s10994-013-5355-6.
- Vladimir Vovk, David Lindsay, Ilia Nourtdinov, and Alexander Gammernan. Mondrian confidence machine. On-line Compression Modelling Project Working Paper 4, Royal Holloway, University of London, 2003.
- Vladimir Vovk, Alexander Gammernan, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, New York, 2005. doi: 10.1007/b106715.
- Nien-Shao Wang, Duygu Nur Yaldiz, Yavuz Faruk Bakman, and Sai Praneeth Karimireddy. Conformal prediction adaptive to unknown subpopulation shifts, 2025. arXiv:2506.05583.
- Yanfei Zhou, Rizal Fathony, Nam H. Nguyen, and Matteo Sesia. Audited conformal prediction for classification under unknown distribution shift, 2026. arXiv:2606.14909.